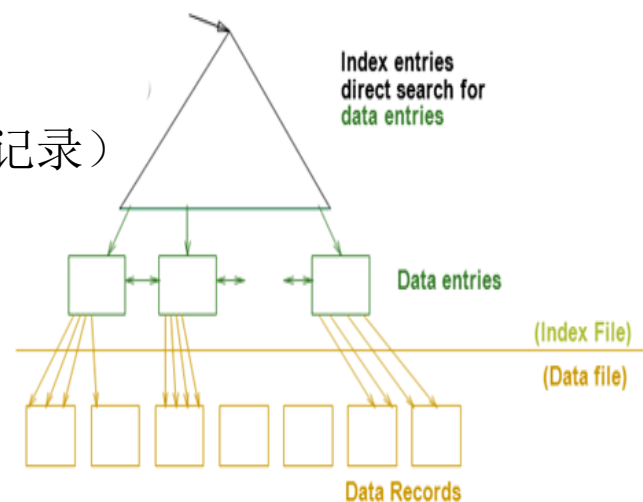


Implementation of Relational Operations

SUN YAT-SEN UNIVERSITY

Review

- **Cost Model**（代价模型）
 - 只关心磁盘块的IO次数
- 文件由页组成，而每个页包含一组记录
 - Record id = <page id, slot #>
- 索引文件由两部份组成
 1. 数据项部分
 - **Data Entry**(数据项) \iff **data record**（数据记录）
 2. 引导部份
 - 树索引技术 $\text{Cost} = \log_F N$ (2~4 I/Os)
 - Hash索引 1~2 I/Os
- 索引的 **clustered**?
- 外排序算法 $\text{COST} = 2N(1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil)$



Introduction

- Next topic: QUERY PROCESSING(查询求值)
- Some database operations are EXPENSIVE
- Huge performance gained by being “smart”
 - We’ll see 10,000x over naïve approach
- Main weapons are:
 - clever implementation techniques for operators
 - Exploiting(利用) relational algebra “equivalences”
 - using statistics and cost models to choose

Simple SQL Refresher

- *SELECT <list-of-fields>*
FROM <list-of-tables>
WHERE <condition>

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='A'
```

A Really **Bad** Query Optimizer

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='A'
```


■ For each **Select-From-Where** query block

Create a **plan** (执行计划) that:

- Forms the cross product of the FROM clause
- Applies the WHERE clause

■ (Then, as needed:

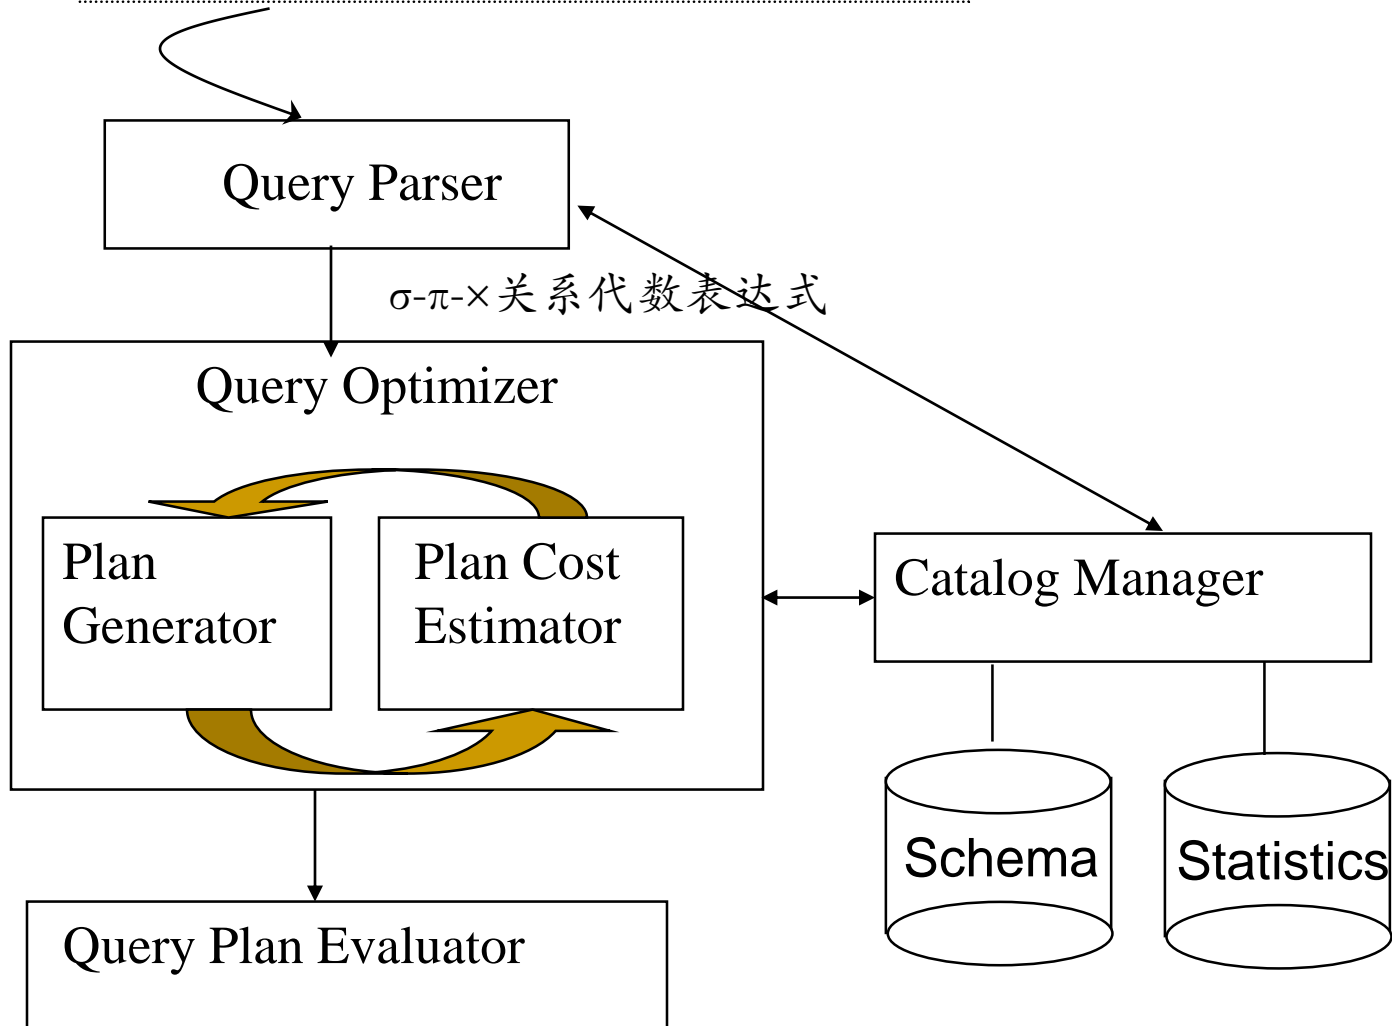
- Apply the GROUP BY clause
- Apply the HAVING clause
- Apply any projections and output expressions
- Apply duplicate elimination and/or ORDER BY)



Cost-based Query Sub-System


Queries

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='A'
```



The Query Optimization Game

- Goal is to pick a “good” plan
 - Good = low expected cost, under *cost model*
 - Degrees of freedom:
 - access methods
 - physical operators
 - operator orders



- Roadmap for this topic:
 - *First:* implementing individual operators
 - *Then:* optimizing multiple operators

Relational Operations

- We will consider how to implement:
 - Selection (σ) Select a subset of rows.
 - Projection (π) Remove unwanted columns.
 - Join (\bowtie) Combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
- Q: What about Intersection?

$$R \cap S = R - (R - S)$$

Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

■ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
- $[S]=500$, $P_S=80$.

■ Reserves:

- Each tuple is 40 bytes, 100 tuples per page, 1000 pages.
- $[R]=1000$, $P_R=100$.

Simple Selections

```
SELECT *  
FROM Reserves R  
WHERE R.rname < 'C%'
```

$$\sigma_{R.attr op value}(R)$$

- How best to perform? Depends on:
 - what indexes are available
 - expected size of result
- **Size of result** approximated as
*(size of R) * selectivity*
 - *Selectivity*(选择性、缩减因子、选中率) estimated via statistics – we will discuss shortly.

Our options ...

$$\sigma_{R.attr op value}(R)$$

- If no appropriate index exists:

Must scan the whole relation

cost = [R].

For “reserves” = 1000 I/Os.

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname < 'C%'
```

```
[S]=500, pS=80.  
[R]=1000, PR=100
```

Our options ...

$$\sigma_{R.attr \text{ op value}}(R)$$

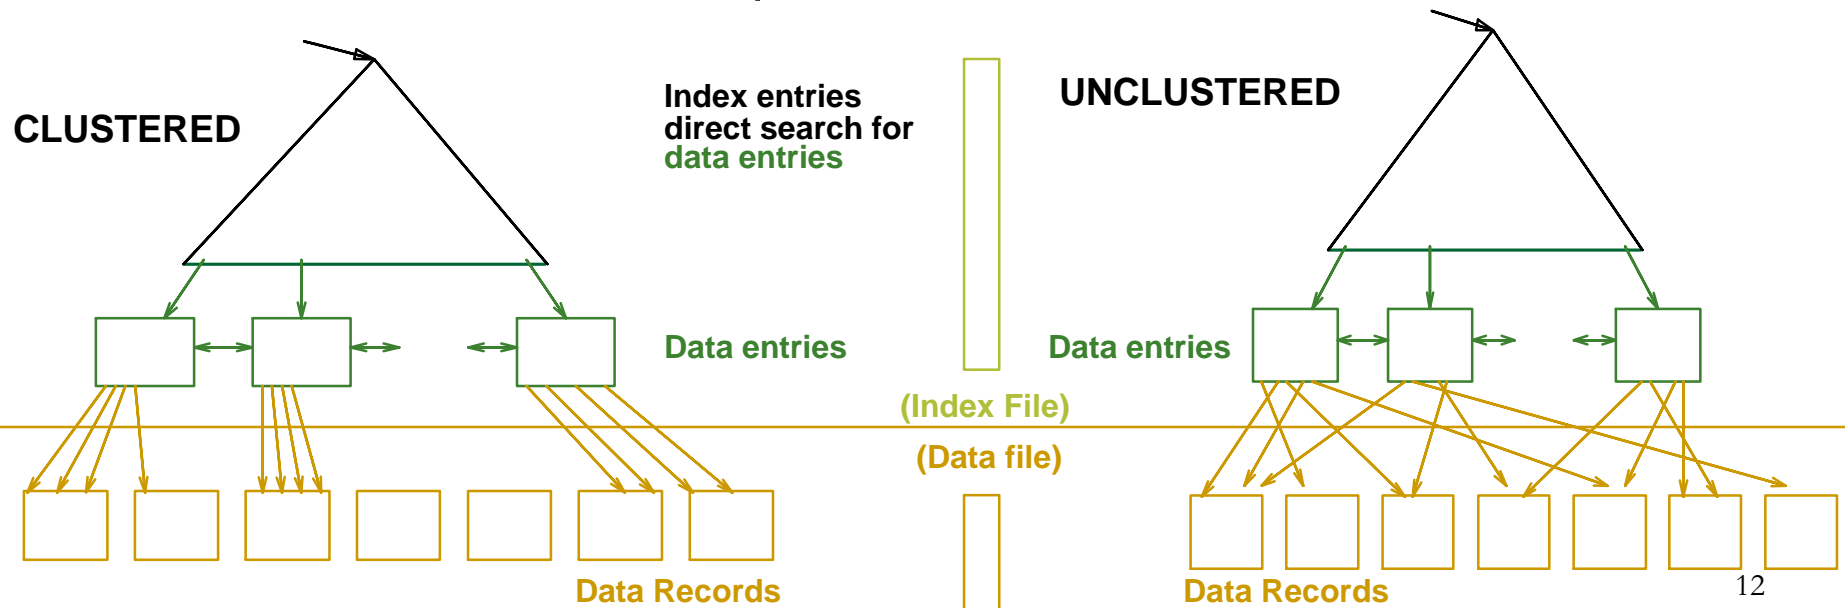
```
SELECT *
FROM   Reserves R
WHERE  R.rname < 'C%'
```

- With index on selection attribute:
 1. Use index to find qualifying **data entries**
 2. Retrieve corresponding **data records**

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

Total cost = cost of step 1 + cost of step 2

- For “Reserves”, if selectivity = 10% (100 pages, 100*100 tuples):
 - If **clustered** index, cost is a little over **100** I/Os;
 - If **unclustered**, could be up to **100*100** I/Os! ... unless ...



Refinement for unclustered indexes


```
SELECT *  
FROM Reserves R  
WHERE R.rname < 'C%'
```

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

1. Find qualifying data entries.
2. Sort the **rids** of the data records to be retrieved.
3. Fetch rids in order.

100*100 tuples

Each data page is looked at just once (though # of such pages likely to be higher than with clustering).



General Selection Conditions

(day<8/9/94 AND rname='Paul') OR bid=5 OR sid=3

- First, convert to conjunctive normal form (CNF)-合取范式:
 - *(day<8/9/94 OR bid=5 OR sid=3) AND (rname='Paul' OR bid=5 OR sid=3)*
- We only discuss the case with no ORs
- Terminology – 索引匹配条件项:
 - A **B-tree** index matches terms(条件项) that involve only attributes in a *prefix* of the search key. e.g.:
 - Index on *<a, b, c>* matches *a=5 AND b= 3*, but not *b=3*.

2 Approaches to General Selections

$$\sigma_{R.attrOpvalueand...}(R)$$

Approach I:

1. Find the *cheapest access path*(访问路径)
 2. retrieve tuples using it
 3. Apply **any remaining terms** that **don't match** the index
-
- ▣ *Cheapest access path*: An index or file scan that we estimate will require the fewest page I/Os.

Cheapest Access Path - Example

query: $\sigma_{\text{day} < 8/9/94 \text{ AND } \text{bid}=5 \text{ AND } \text{sid}=3}(R)$

some options:

B+tree index on day; check bid=5 and sid=3 afterward.

hash index on <bid, sid>; check day<8/9/94 afterward.

- *How about a B+tree on <rname, day>?*
- *How about a B+tree on <day, rname>?*
- *How about a Hash index on <day, rname>?*

2 Approaches to General Selections (Contd.)

Approach II: use 2 or more matching indexes.

1. From each index, get set of rids
2. Compute intersection of rid sets
3. Retrieve records for rids in intersection
4. Apply any remaining terms

EXAMPLE: $\sigma_{\text{day} < 8/9/94 \text{ AND } \text{bid}=5 \text{ AND } \text{sid}=3}(R)$

Suppose we have an index on *day*, and another index on *sid*.

- ❑ Get rids of records satisfying *day*<8/9/94.
- ❑ Also get rids of records satisfying *sid*=3.
- ❑ Find intersection, then retrieve records, then check *bid*=5.

Projection

```
SELECT  DISTINCT  
        R.sid, R.bid  
FROM    Reserves R
```

```
[S]=500,  pS=80.  
[R]=1000, PR=100
```

- Issue is removing **duplicates**.
- Use sorting!!
 1. Scan R, extract only the needed attributes
 2. Sort the resulting set
 3. Remove adjacent duplicates

Cost:

- writes to temp table at each step!

Reserves with size ratio **0.25** = 250 pages.

With **20 buffer pages** can sort in **2** passes, so:

$$1000 + 250 + 2 * 2 * 250 + 250 = 2500 \text{ I/Os}$$

$$2N(1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil)$$

Projection -- improved

1. Scan R, extract only the needed attributes
2. Sort the resulting set
3. Remove adjacent duplicates

```
SELECT  DISTINCT  
        R.sid, R.bid  
FROM    Reserves R
```

- Avoid the temp files, work on the fly(实时流水线):
 - Modify Pass 0 of sort to eliminate unwanted fields.
 - Modify Pass 1+ to eliminate duplicates.

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

Cost:

Reserves with size ratio $0.25 = 250$ pages.

With 20 buffer pages can sort in 2 passes, so:

1. Read 1000 pages
2. Write 250 (in runs of 40 pages each) = 7 runs
3. Read and merge runs (20 buffers, so 1 merge pass!)

Total cost = 1000 + 250 + 250 = 1500.

$$1000 + 250 + 2 * 2 * 250 + 250 = 2500 \text{ I/Os} \quad \left(1 + \left\lceil \log_{B-1} \left\lceil N / 2B \right\rceil \right\rceil\right)$$

Other Projection Tricks

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

If an index search key contains all wanted attrs:



- Do *index-only* (唯索引) scan
 - Apply projection techniques to data entries (*much smaller!*)

If a B+Tree index search key *prefix* has all wanted attrs:

- Do *in-order index-only* (有序唯索引) scan
 - Compare adjacent tuples on the fly (*no sorting required!*)

N

Joins

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- Joins are very common.
- $R \times S$ is large; so, $R \times S$ followed by a selection is inefficient.
- Many approaches to reduce join cost.
 - Join techniques we will cover today:
 1. Nested-loops join
 2. Index-nested loops join
 3. Sort-merge join
 4. Hash-join

Simple Nested Loops Join (简单的嵌套循环连接算法)

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

```
R × S:  foreach tuple r in R do
        foreach tuple s in S do
            if r_i == s_j then add <r, s> to result
```

$$\text{Cost} = [R] + (P_R * [R]) * [S] = 1000 + 100 * 1000 * 500 \text{ IOs}$$

- At 10ms/IO, Total time: ??? = 1000 + 50,000,000 IOs
140小时

- What if smaller relation (S) was “outer”? $[S] + (P_s * [S]) * [R]$
500 + 80*500*1000
- What assumptions are being made here? **Not all in memory**
- What is cost if one relation can fit entirely in memory? $[R] + [S]$

Page-Oriented Nested Loops Join

(页嵌套循环连接算法)

$R \bowtie S$: foreach page b_R in R do
 foreach page b_S in S do
 foreach tuple r in b_R do
 foreach tuple s in b_S do
 if $r_i == s_j$ then add $\langle r, s \rangle$ to result

$[S]=500, \quad p_S=80.$ $[R]=1000, \quad P_R=100$

$$\text{Cost} = [R] + [R] * [S] = 1000 + 1000 * 500 = 501,000$$

1小时

- If smaller relation (S) is outer, cost = $500 + 500 * 1000 = 500,500$

1%

- Much better than naïve per-tuple approach!

Block Nested Loops Join


```
foreach block  $b_R$  in R do  
  foreach page  $b_S$  in S do
```

...

(块嵌套循环连接算法)

$$\text{Cost} = [R] + [R] * [S]$$

- Page-oriented NL doesn't exploit extra buffers :(
- Idea to use memory efficiently:




Cost: Scan outer + (#outer blocks * scan inner)

$$\#outer\ blocks = \left\lceil \#of\ pages\ of\ outer / blocksize \right\rceil$$

$$\text{Cost} = [R] + [R]/N * [S]$$

Examples of Block Nested I



- Say we have $B = 100+2$ memory buffers
- Join cost = [outer] + (#outer blocks * [inner])
 $\#outer\ blocks = [outer] / 100$
- With R as outer ($[R] = 1000$):

$[S]=500, p_S=80.$
 $[R]=1000, P_R=100$

 - Scanning R costs 1000 IO's (*done in 10 blocks*)
 - Per block of R, we scan S; costs $10*500$ I/Os
 - Total = $1000 + 10*500=6000$. **Cost = $[R] + [R]/N * [S]$**
- With S as outer ($[S] = 500$):
 - Scanning S costs 500 IO's (*done in 5 blocks*)
 - Per block of S, we scan R; costs $5*1000$ IO's
 - Total = $500 + 5*1000=5500$.

55秒


1/10000

Index Nested Loops Join

(索引嵌套循环连接算法)

$R \bowtie S$: foreach tuple r in R do

foreach tuple s in S where $r_i == s_j$ do
add $\langle r, s \rangle$ to result



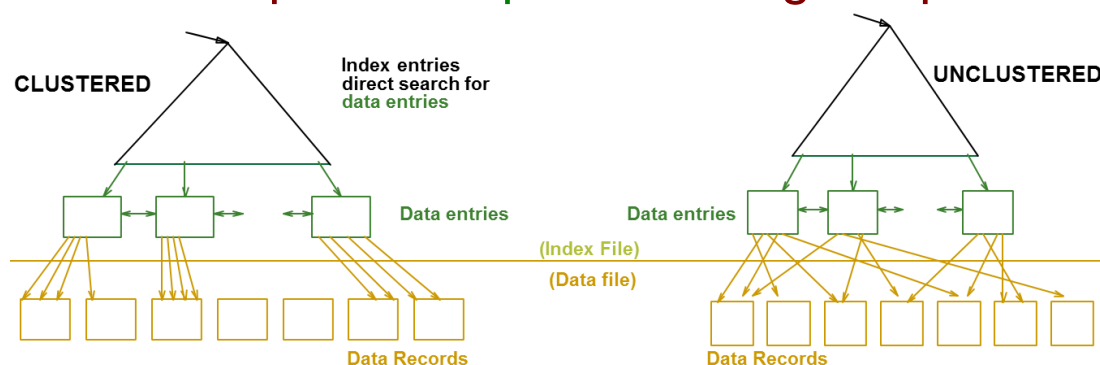
Index Nested Loops Join

[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

$R \bowtie S$: foreach tuple r in R do
 foreach tuple s in S where $r_i == s_j$ do
 add $\langle r, s \rangle$ to result

Cost = $[R] + ([R] * P_R) * \underline{\text{cost to find matching } S \text{ tuples}}$

- If index uses **Alt. 1**, cost = cost to traverse tree from root to leaf.
- For **Alt. 2 or 3**:
 1. Cost to lookup RID(s); typically 2-4 IO's for B+Tree.
 2. Cost to retrieve records from RID(s); depends on clustering.
 - ❑ **Clustered index**: 1 I/O per page of matching S tuples.
 - ❑ **Unclustered**: up to 1 I/O per matching S tuple.



Sort-Merge Join

(排序归并连接算法)

Example:

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

1. Sort R on join attr(s)
2. Sort S on join attr(s)
3. Scan sorted-R and sorted-S
in tandem, to find matches

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Cost of Sort-Merge Join

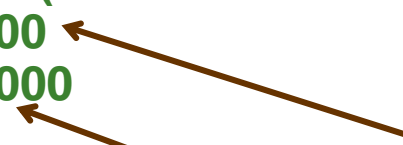
[S]=500, $p_S=80$.
[R]=1000, $P_R=100$

- Cost: Sort R + Sort S + ([R]+[S])
 - But in worst case, last term could be [R]*[S] (*very unlikely!*)
 - Q: what is worst case? 每个元组的连接属性值都一样

Suppose $B = 35$ buffer pages: $2N(1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil)$

- Both R and S can be sorted in 2 passes
- Total join cost = $4*1000 + 4*500 + (1000 + 500) = 7500$

Suppose $B = 300$ buffer pages:

- Again, both R and S sorted in 2 passes
 - Total join cost = 7500
- Join cost = [outer] + (#outer blocks * [inner])
 $500 + 500/(35-2) * 1000$
 $500 + 500/(300-2) * 1000$
- 

Block-Nested-Loop cost = 2500 ... 15,000

Other Considerations ...

$[S]=500, \quad p_S=80.$ $[R]=1000, \quad P_R=100$

1. An important refinement:

Do the join during the final merging pass of sort !

- If have enough memory, can do:
 1. Read R and write out sorted runs
 2. Read S and write out sorted runs
 3. Merge R-runs and S-runs, while finding $R \bowtie S$ matches

$$\text{Cost} = 3*[R] + 3*[S]$$

Q: how much memory is “enough” ?

$$B > 2 * \sqrt{\text{Max}([R], [S])}$$

超过R和S的有序段的数目

2. Sort-merge join an especially good choice if:

- one or both inputs are **already sorted** on join attribute(s)
- output is **required to be sorted** on join attributes(s)

$$2N(1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil)$$

Hash-join(哈希连接算法)

$$\text{Cost} = 3 * [R] + 3 * [S]$$

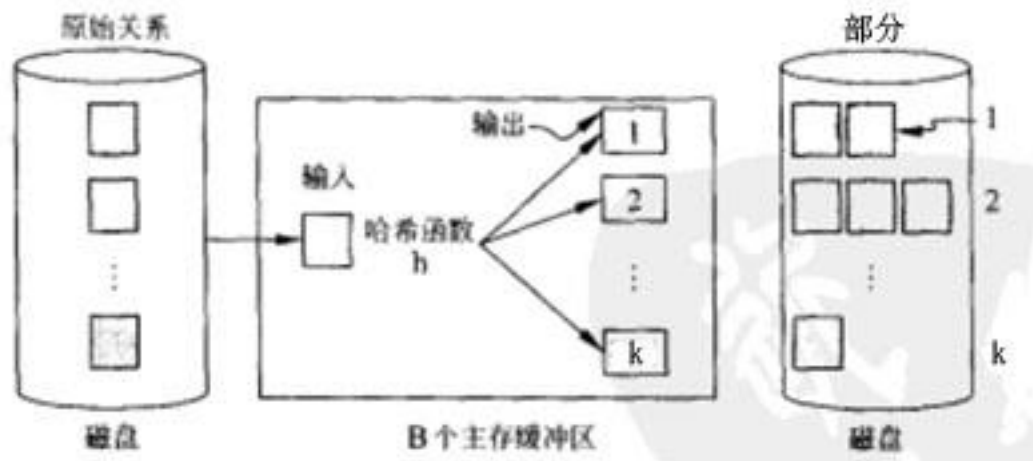
- **基本思想**：先（第1步）对两个关系的连接属性应用同一个 Hash 函数，将每个关系划分到 k 个部分。显然可以保证部分 i 的 R_i 元组只能与同一部分的 S_i 元组连接。
然后（第2步），我们可以读入较小关系 R 的一个部分 R_i （即将 R_i 全部读入内存中），然后扫描关系 S 的对应部分 S_i ，来找到匹配的元组。
- 算法分为**两个阶段**
 - partitioning phase划分阶段
 - probing phase连接阶段 (也称匹配、探查阶段)

Hash-join(哈希连接算法)-划分阶段

$$\text{Cost} = 3 * [R] + 3 * [S]$$

■ partitioning phase 划分阶段

- 使用同一个Hash函数 h 对关系 R 和 S 进行划分。完成后，所有元组划分为 k 个部分，每个部分中的元组有相同的哈希值（对元组的连接属性计算哈希函数 h ）。



- 在划分阶段需要扫描关系 R 和 S ，还要写出一次，因此

$$\text{Cost} = 2 * [R] + 2 * [S]$$

Hash-join(哈希连接算法)-连接阶段

$$\text{Cost} = 3 * [R] + 3 * [S]$$

■ probing phase连接阶段

- 比较属于同一部分的 R_i 元组和 S_i 元组，找到匹配元组。
- 实际处理时，为了减少CPU开销，使用与 h 不同的Hash函数 h_2 对较小的部分（如 R_i ）进行划分，在内存中创建 R_i 部分的哈希表。

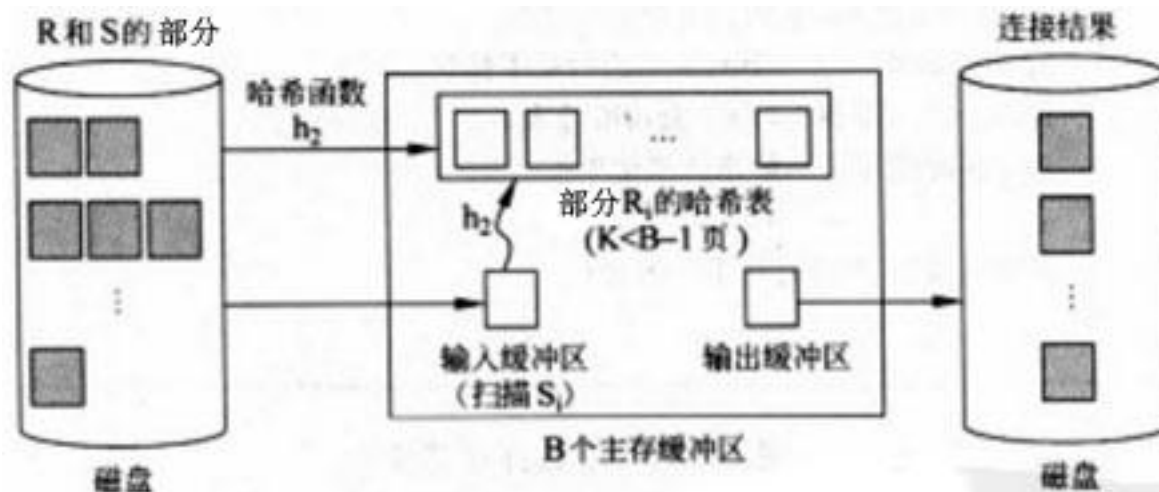


图 14.11 哈希连接算法的连接阶段

- 在连接阶段，每个部分需要扫描一次，故 $\text{Cost} = [R] + [S]$

Summary

- A virtue of relational DBMSs:

queries are composed of a few basic operators

- Many alternative implementation techniques for each operator
 - No universally superior technique for most operators.
- Must consider available alternatives
 - Called “Query optimization” -- we will study this topic soon!
- 要求:
 - 根据给定条件, 计算关系运算 σ 、 π 和 \bowtie 的 COST, 如:

$$\text{Cost} = [R] + [R] * [S]$$

$$\text{Cost} = [R] + [R]/N * [S]$$

$$\text{Cost: Sort R + Sort S} + ([R]+[S])$$