


# Interconnection Network

# Interconnection Networks[互连网络]

- An **Interconnection Network (ICN)** is a **programmable system** that transports data between terminals
  - To **hold our parallel machines together**
    - at the core of parallel computer architecture
  - **Share basic concept** with LAN/WAN
  - but **very different trade-offs** due to very different time scale/requirements



# Interconnection Networks[互连网络]

- **Interconnection networks can be grouped into four domains**
  - Depending on number and proximity of devices to be
    - **Wide-Area Networks**[广域网]
    - **Local-Area Networks**[局域网]
    - **System-Area Networks**[系统区域网络]
    - **On-Chip Networks**[片上网络]

# Wide-Area Networks

- **Interconnect systems distributed across the globe**
  - Internetworking support is required
  - Millions of devices interconnected
- **Maximum interconnect distance**
  - many thousands of kilometers

# Local-Area Networks

- **Interconnect autonomous computer systems**
  - Machine room or throughout a building or campus
  - Hundreds of devices interconnected
- **Maximum interconnect distance**
  - Few meters to tens of kilometers
  - Example (most popular):
    - Ethernet, with 10 Gbps over 40Km

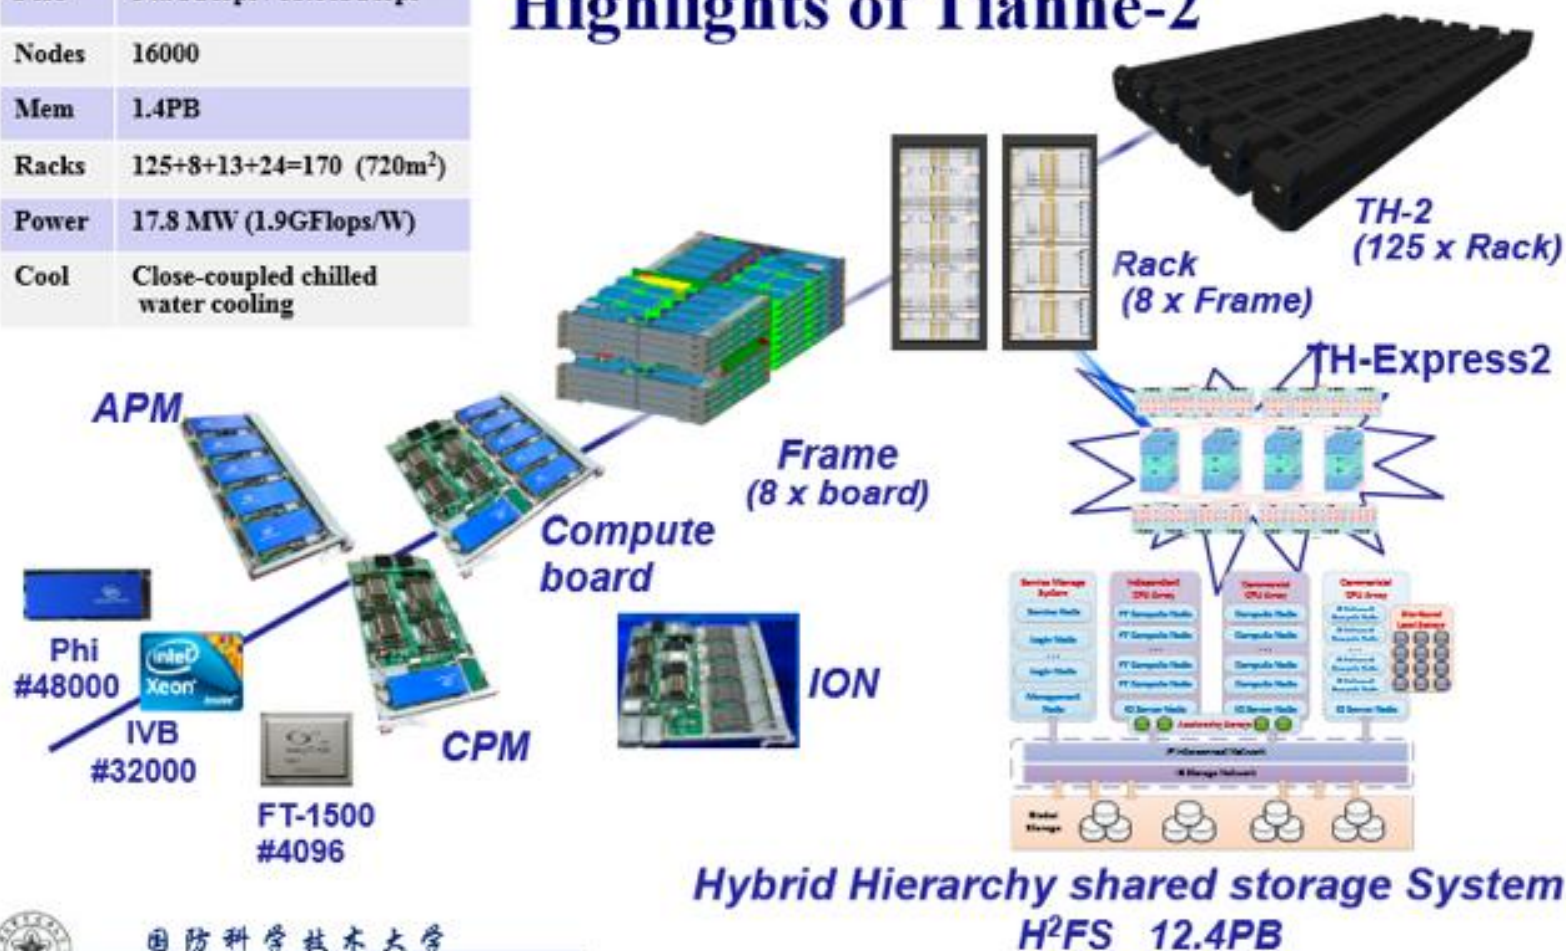
# System-Area Networks

- **Interconnects within one “machine”**
  - Interconnect in a multi-processor system
  - Interconnect in a supercomputer
  - Hundreds to thousands of devices interconnected
    - Tianhe-2 supercomputer (16K nodes)
- **Maximum interconnect distance**
  - Fraction to tens of meters (typical)
  - A few hundred meters (some)
    - InfiniBand: 120 Gbps over a distance of 300m

# System-Area Networks

Perf	54.9PFlops / 33.86PFlops
Nodes	16000
Mem	1.4PB
Racks	125+8+13+24=170 (720m <sup>2</sup> )
Power	17.8 MW (1.9GFlops/W)
Cool	Close-coupled chilled water cooling

## Highlights of Tianhe-2



国防科学技术大学  
 National University of Defense Technology

# On-Chip Networks

- **Interconnect** within a single chip
  - Devices are micro-architectural elements
    - Caches, directories, processor cores
  - Currently, designs with tens of devices are common
    - Ex: IBM Cell, Intel multicores, Tile processors
- **Proximity:** millimeters

**We are concerned with On-Chip and System-Area Networks**



# Why Study Interconnects?

- Provide **external connectivity** from system to outside world
  - Also, connectivity within a single computer system at many levels
- Interconnection networks should be **well designed**
  - To transfer the **maximum amount of information**
  - Within the **least amount of time** (and cost, power constraints)
- **Application:**
  - managing communication can be critical to performance

# Why Study Interconnects? (cont.)

- **Trends:** high demand on communication bandwidth
  - increased computing power and storage capacity
  - switched networks are replacing buses
- **Computer architects/engineers must understand interconnect problems and solutions**
  - in order to effectively design and evaluate systems

# Interconnection network


- **Characteristics of communication**

- Latency must be **low**
- Bandwidth must be **high**
- Message characteristics
  - Many **small messages** with several fixed sizes
  - The largest message is one cache block size (64 or 128 bytes)

- **Implications**

- **Only two layers are sufficient:** link level and node level
- Communication protocol must be **simple**
  - no packet dropping, no elaborate flow control


# Nodes and Links



- **A node encapsulates**
  - one or more processors
  - communication controller/assist (CA)
  - router
  - Routers are connected by links
  - Links may be unidirectional or bidirectional

# Link and Channel

- **Link** = **set of wires interconnecting a pair of nodes**
  - Link can be **unidirectional** or **bidirectional**
  - Unit of transfer determined by **flow control protocol** (flit)
    - **Flit size** determined by **link latency** and amount of **buffering** available
- **Channel** = **link + sender + receiver**
- A message transmitted over a channel is broken into **packets**
  - A packet has a header, trailer, and payload



# Design Considerations

- **Application requirements**
  - Number of terminals or ports to support
  - Peak bandwidth of each terminal
  - Average bandwidth of each terminal
  - Latency requirements
  - Message size distribution
  - Expected traffic patterns
  - Required quality of service
  - Required reliability and availability

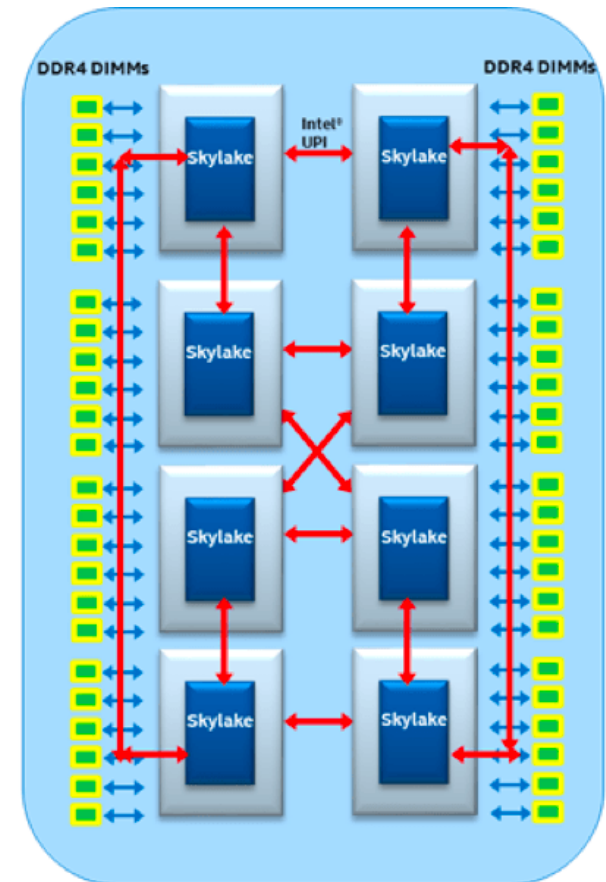
# Design Considerations

- **Job of an interconnection network:**
  - to transfer information from source node to dest. node
  - support network transactions that realize application
- **Latency as small as possible**
  - As many concurrent transfers as possible
- **Cost as low as possible**

# Example Requirements

- **Example requirements for a coherent processor-memory interconnect**

- Processor ports 1-2048
- Memory ports 1-4096
- Peak BW 8 GB/s
- Average BW 400 MB/s
- Message Latency 100 ns
- Message size 64 or 576 bits
- Traffic pattern arbitrary
- Quality of service none
- Reliability no message loss
- Availability 0.999 to 0.99999

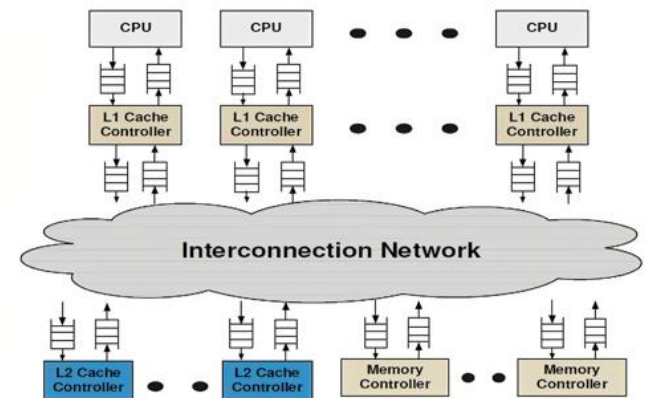




# Technology constraints

- **Technology constraints**


- Signaling rate
- Chip pin count (if off-chip networking)
- Area constraints (typically for on-chip networking)
- Chip cost
- Circuit board cost (if backplane boards needed)
- Signals per circuit board
- Signals per cable
- Cable cost
- Cable length
- Channel and switch power constraints



# Off-chip vs. On-chip ICNs


- **Off-chip: I/O bottlenecks**

- Pin-limited bandwidth
- Inherent overheads of off-chip I/O transmission



- **On-chip**

- Wiring constraints
  - Horizontal and vertical layout
  - Short, fixed length
- Power
  - Consume 10-15% or more of die power budget
- Latency
  - Different order of magnitude



# Main Aspects of an ICN

- **Topology**
  - Static arrangement of channels and nodes in a network
- **Routing**
  - Determines the set of paths a packet can follow
- **Flow control[流控]**
  - Allocating network resources (channels, buffers, etc.) to packets and managing contention

# Main Aspects of an ICN

- **Switch microarchitecture**
  - Internal architecture of a network switch
- **Network interface**
  - How to interface a terminal with a switch
- **Link architecture**
  - Signaling technology and data representation on the channel

# I. Network Topology

# Network Topology


- **Network topology = shape of the network**
  - Determines the distance that a message needs to travel (in # links or network hops) from one node to another
- **Important characteristics**
  - **Diameter**: largest distance
  - **Average distance**
  - **Degree**: how many links are connected to a switch

# Example

- For the shown network, compute diameter, avg distance, bisection bandwidth (assuming each link is bidirectional)

**Diameter = 3**

e.g. the distance between node 1 and 6



**Avg distance:**

$d(1,2)=1$ ,  $d(1,3)=2$ ,  $d(1,4)=1$ ,  $d(1,5)=2$ ,  
 $d(1,6)=3$

$d(2,3)=1$ ,  $d(2,4)=2$ ,  $d(2,5)=2$ ,  $d(2,6)=2$

$d(3,4)=2$ ,  $d(3,5)=1$ ,  $d(3,6)=1$


$d(4,5)=1$ ,  $d(4,6)=2$

$d(5,6)=1$

Sum of all distances = 24, number of  
node pairs = 15

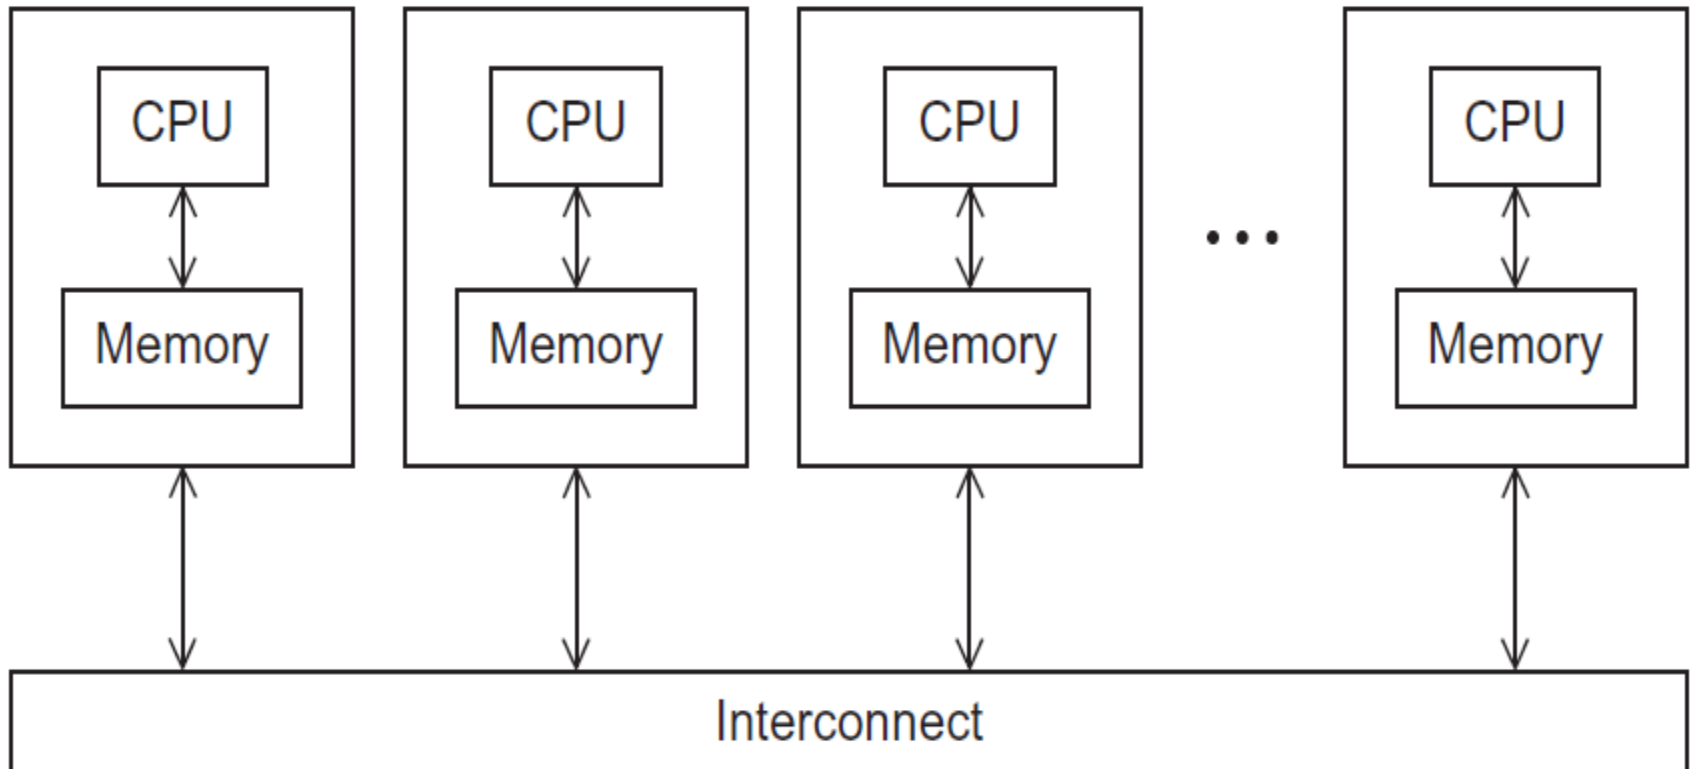
Average distance =  $24/15$

# Recall: Shared Memory System





# Recall: Distributed Memory System



# Interconnection networks

- **Affects performance** of both distributed and shared memory systems.
- Two categories:
  - Shared memory interconnects
  - Distributed memory interconnects

# I.I Shared memory interconnects

# Shared memory interconnects

- **Bus interconnect**

- A collection of parallel communication wires together with some hardware that controls access to the bus.
- Communication wires are **shared by the devices** that are connected to it.
- As the number of devices connected to the bus increases, **contention** for use of the bus **increases**, and performance decreases.

# Shared memory interconnects

- **Switched interconnect**


- Uses **switches** to control the **routing** of data among the connected devices.

- **Crossbar** –

- Allows **simultaneous communication** among different devices.
      - **Faster** than buses.
      - But the **cost** of the switches and links is relatively **high**.


(a)

A crossbar switch connecting 4 processors ( $P_i$ ) and 4 memory modules ( $M_j$ )

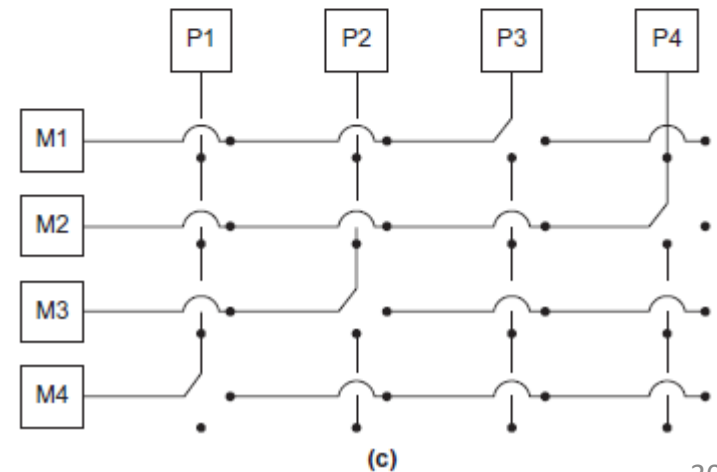


(b)

Configuration of internal switches in a crossbar



(c) Simultaneous memory accesses by the processors



## 1.2 Distributed memory interconnects

# Distributed memory interconnects

- **Two groups**

- **Direct interconnect**

- Each switch is **directly connected** to a processor memory pair, and the switches are connected to each other.


- **Indirect interconnect**

- Switches may **not be directly connected** to a processor.




# ( I ) Direct interconnect

# Direct interconnect



(a)

ring



(b)

toroidal mesh


(二维环面网格)

# Bisection width（等分宽度）


- A measure of “**number of simultaneous communications**” or “**connectivity**”.
- How many simultaneous **communications** can take place “**across the divide**” between the halves?



# Two bisections of a ring



(a)



(b)


# Definitions

- Bandwidth
  - The rate at which a link can transmit data.
    - Usually given in megabits or megabytes per second.
- Bisection bandwidth
  - A measure of network quality.
  - Instead of counting the number of links joining the halves, it sums the bandwidth of the links.

# Fully connected network

- Each switch is directly connected to every other switch.

*impractical*




bisection width =  $p^2/4$

( $p$ : # of nodes)

# Hypercube


- **Highly connected** direct interconnect.
- **Built inductively:**
  - A **one-dimensional hypercube** is a fully-connected system with two processors.
  - A **two-dimensional hypercube** is built from two one-dimensional hypercubes by **joining** “corresponding” switches.
  - Similarly a **three-dimensional hypercube** is built from two two-dimensional hypercubes.

# Hypercubes




(a)

one-



(b)

two-



(c)

three-dimensional

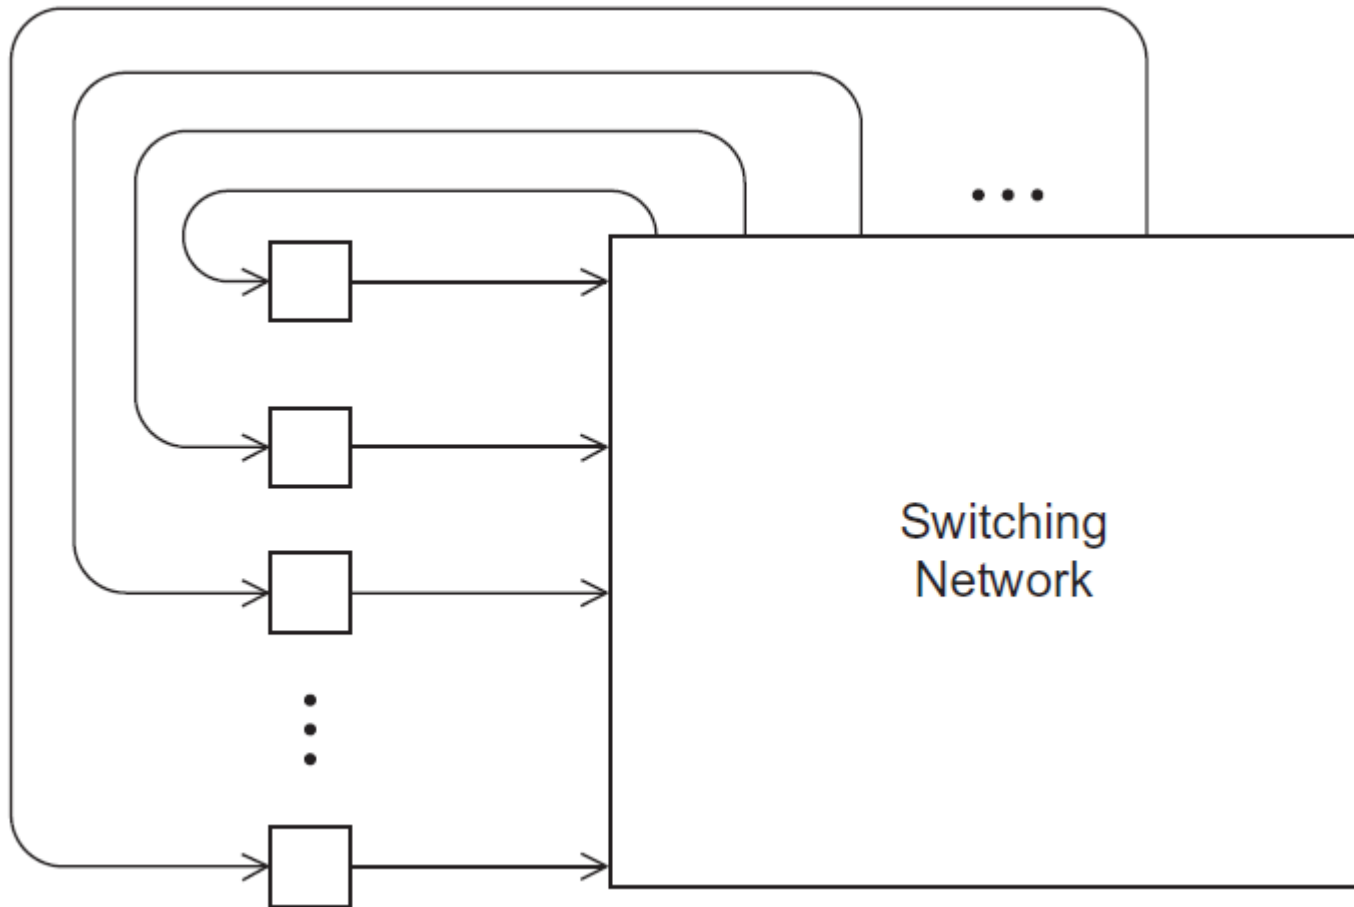


## (2) Indirect interconnect


# Indirect interconnects

- Simple examples of indirect networks:
  - Crossbar
  - Omega network
- Often shown with unidirectional links and a collection of processors
- Each of which has an **outgoing** and an **incoming link**, and a **switching network**.


# A generic indirect network




# Crossbar interconnect for distributed memory



# An omega network



# A switch in an omega network



# More definitions

- Any time data is transmitted, we're interested in **how long it will take** for the data to reach its destination.
- **Latency**
  - The time that elapses between the source's beginning to transmit the data and the destination's starting to receive the first byte.
- **Bandwidth**
  - The **rate** at which the destination receives data after it has started to receive the first byte.

Message transmission time =  $1 + n / b$

latency (seconds)

length of message (bytes)

bandwidth (bytes per second)

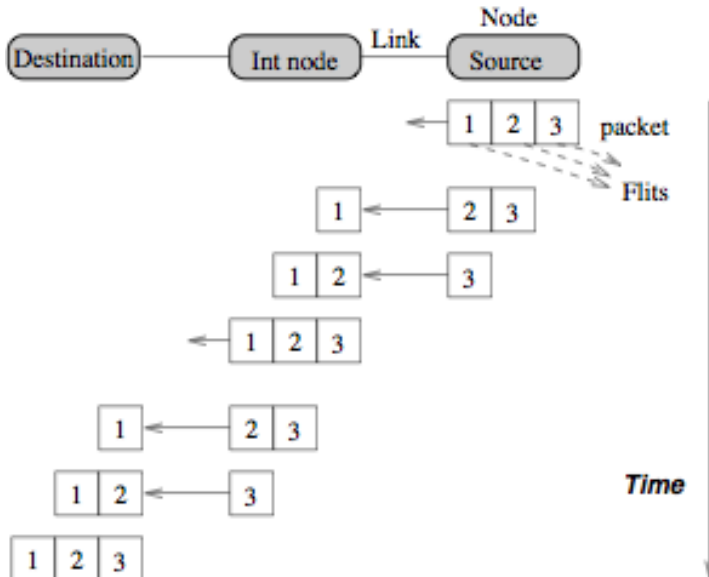


## 2. Routing Policy


# Routing policy

How to send a packet over the network?

**Store and forward:**



**Cut-through:**



- **Store & Forward:**

- Send all flits over a link, then over another, and so on
- $T = h * (T_{\text{xmit}} + T_{\text{switch}})$

- **Cut-through:**

- Send flits in pipeline over the links
- $T = h * T_{\text{switch}} + T_{\text{xmit}}$


- **Implication:**

- cut-through routing reduces the importance of topology

# Routing Policy

- **Minimal vs. non-minimal**
  - **Minimal** = each packet travels the least number of hops
- **Deterministic vs. non-deterministic**
  - **Deterministic** = each sender-dest pair uses a single path for all messages
- **Adaptive vs. non-adaptive**
  - **Adaptive** = message can change path as it is being transmitted
- **Deadlock-free vs. Deadlock possible**
  - Deadlock due to several messages mutually waiting for buffer space

# What is Deadlock?

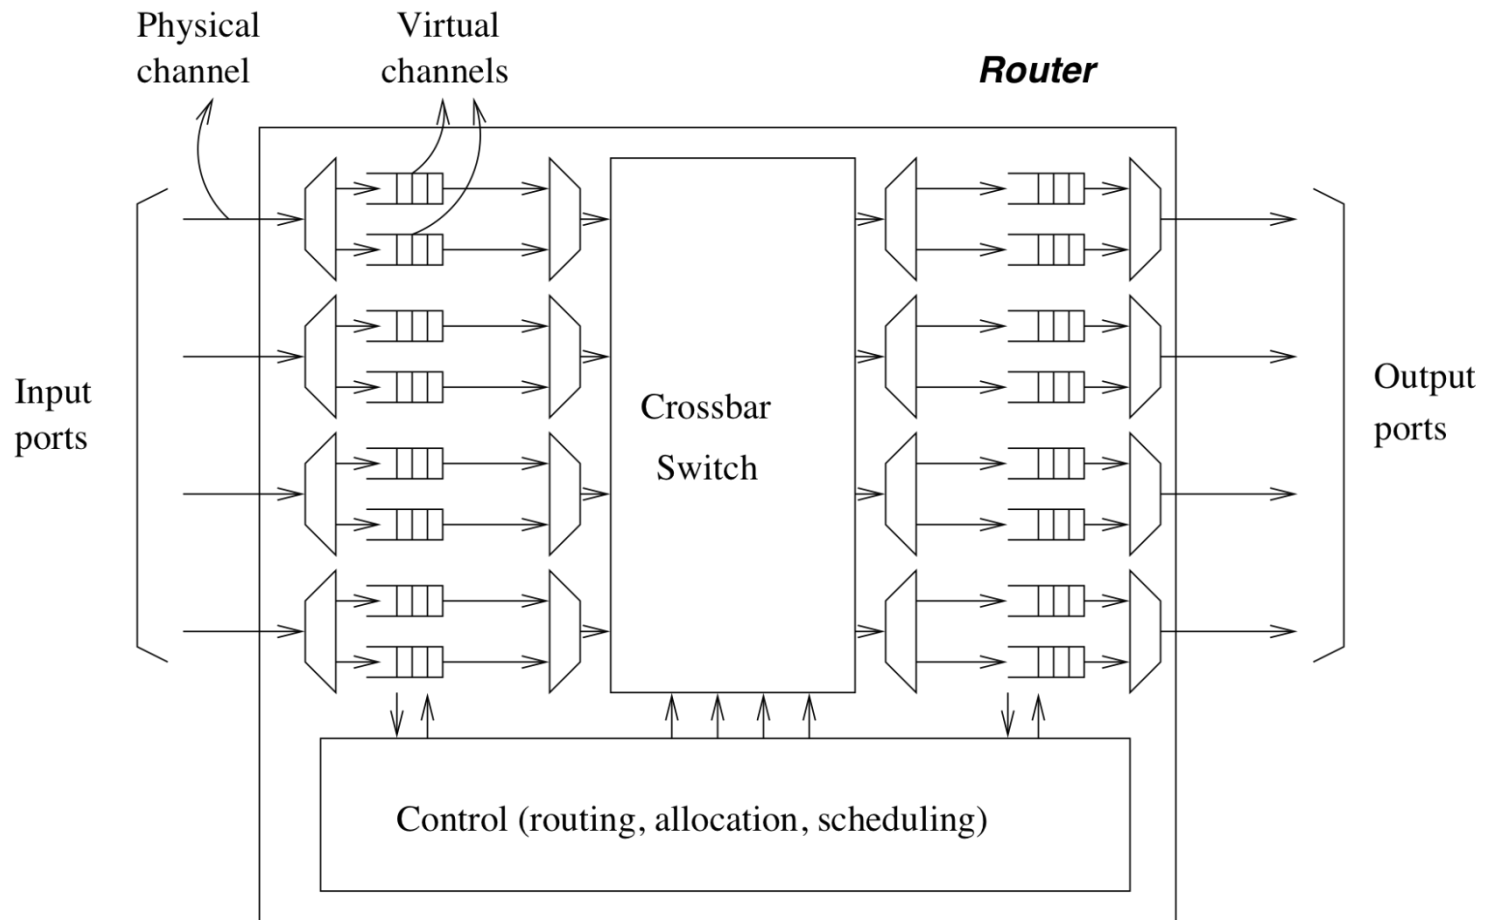


The figure shows four packets filling up the input buffer and need the next output buffer, but the output buffer is already full

- **Inability** of the network to forward packets
  - Due to limited buffer space and cyclic dependence on buffer acquisition
- **Handling a deadlock**
  - **Detect and break**: expensive
  - **Drop packets**: adverse effect in performance and complicates protocol
  - **Avoid**: restriction on routing

# 3. Router Architecture

# Router Architecture



- Router has four input ports and four output ports
- Each channel has two virtual channels (VCs), with separate buffers

# Router Architecture

- Arriving packet sent to the **appropriate VC** (buffer) based on its **header**
- Control logic arbitrates among **multiple input channels** and allocates them to the input of the crossbar switch
  - It also allocates **output VC** to the output of the switch
  - It **arbitrates multiple input ports** that want to send packet to the same output port

# Router Architecture

- Switch typically implemented using a **crossbar** (to minimize latency and collision)
  - Multiple flits from different VCs can go on the switch as long as they go to different output VCs
- Steps can also be pipelined to improve **router throughput**



# Summary

- **Interconnection network**
  - Basic concepts
  - Design consideration
  - Main aspects
- **Network Topology**
- **Routing Policy**
- **Router Architecture**