



《模式识别》

第二章 特征提取与表示

马锦华

<https://cse.sysu.edu.cn/teacher/MaJinhua>

SUN YAT-SEN University



声明：该PPT只供非商业使用，也不可视为任何出版物。由于历史原因，许多图片尚没有标注出处，如果你知道图片的出处，欢迎告诉我们 majh8@mail.sysu.edu.cn.



课程目录（暂定）

❑ 第一章	课程简介与预备知识	6学时
❑ 第二章	特征提取与表示	6学时
❑ 第三章	主成分分析	3学时
❑ 第四章	归一化、判别分析、人脸识别	3学时
❑ 第五章	EM算法与聚类	3学时
❑ 第六章	贝叶斯决策理论	3学时
❑ 第七章	线性分类器与感知机	3学时
❑ 第八章	支持向量机	3学时
❑ 第九章	神经网络、正则项和优化方法	3学时
❑ 第十章	卷积神经网络及经典框架	3学时
❑ 第十一章	循环神经网络	3学时
❑ 第十二章	Transformer	3学时
❑ 第十三章	自监督与半监督学习	3学时
❑ 第十四章	开放世界模式识别	6学时



第一部分：模型拟合与局部特征

What we will learn today?

- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector



What we will learn today?



- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector

Fitting as Search in Parametric Space



- **Let's say we have chosen a parametric model for a set of features:**
 - For example, we have a line equation that we want to fit to a set of edge points
- **We can 'search' in parameter space by trying many potential parameter values and see which set of parameters 'agree'/fit with our set of features**
- **Three main questions:**
 1. What model represents this set of features best?
 2. Which of several model instances gets which feature?
 3. How many model instances are there?
- **Computational complexity is important**
 - It is infeasible to examine every possible set of parameters and every possible combination of features

Example: Line Fitting

- Why fit lines? Many objects characterized by presence of straight lines:




Difficulty of Line Fitting



- **Extra edge points (clutter), multiple models:**

- Which points go with which line, if any?



- **Only some parts of each line detected, and some parts are missing:**

- How to find a line that bridges missing evidence?

- **Noise in measured edge points, orientations:**

- How to detect true underlying parameters?

Voting as a fitting technique


- **It's not feasible to check all combinations of features by fitting a model to each possible subset. For example, the naïve line fitting between every pair of two points is $O(N^2)$.**
- **Voting is a general technique where we let the features vote for all models that are compatible with it.**
 - Cycle through features, cast votes for model parameters.
 - Look for model parameters that receive a lot of votes.
- **Noise & clutter features will cast votes too, but typically their votes should be inconsistent with the majority of “good” features.**
- **Ok if some features not observed, as model can span multiple fragments.**

RANSAC (RANdom SAmple Consensus)

- **RANSAC [Fischler & Bolles 1981]**
 - Approach: we want to avoid the impact of outliers, so let's look for “inliers”, and use only those.
- Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.
- **RANSAC loop:**
 1. Randomly select a seed group of points on which to perform a model estimate (e.g., a group of good points)
 2. Compute model parameters from seed group
 3. Find inliers to this model
 4. If the number of inliers is sufficiently large, re-compute least-squares estimate of model on all of the inliers
 - Keep the model with the largest number of inliers


RANSAC Line Fitting Example

- **Task: Estimate the best line**
 - How many points do we need to estimate the line?



RANSAC Line Fitting Example


- Task: Estimate the best line



Sample two points


RANSAC Line Fitting Example

- **Task: Estimate the best line**




RANSAC Line Fitting Example

- Task: Estimate the best line




RANSAC Line Fitting Example

- Task: Estimate the best line



RANSAC Line Fitting Example


- **Task: Estimate the best line**



Repeat, until we get a good result.

RANSAC Line Fitting Example

- Task: Estimate the best line



Algorithm of RANSAC

Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

Draw a sample of n points from the data
uniformly and at random

Fit to that set of n points

For each data point outside the sample

Test the distance from the point to the line
against t ; if the distance from the point to the line
is less than t , the point is close

end

If there are d or more points close to the line
then there is a good fit. Refit the line using all
these points.

end

Use the best fit from this collection, using the
fitting error as a criterion

RANSAC: How many iterations “k”?

- **How many samples (iterations) are needed?**
 - Suppose w is fraction of inliers (points from line).
 - n points needed to define hypothesis (e.g. 2 for lines)
 - k samples chosen.
- **Prob. that a single sample of n points is correct: w^n**
- **Prob. that a single sample of n points fails: $1 - w^n$**
- **Prob. that all k samples fail is: $(1 - w^n)^k$**
- **Prob. that at least one of the k samples is correct:**
 $1 - (1 - w^n)^k$
- \Rightarrow Choose k high enough to keep this below desired failure rate.



RANSAC: Computed k (p=0.99)

$$k = \frac{\log(z)}{\log(1 - w^n)} \quad z = 1 - p$$

Sample size n	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

Refining RANSAC estimate

- RANSAC computes its best estimate from a minimal sample of n points, and divides all data points into inliers and outliers using this estimate.
- We can improve this initial estimate by estimation over all inliers (e.g. with standard least-squares minimization).
- But this may change inliers, so alternate fitting with re-classification as inlier/outlier.



RANSAC: Pros and Cons

- **Pros:**

- General method suited for a wide range of model fitting problems
- Easy to implement and easy to calculate its failure rate ($1-p$)

- **Cons:**

- Only handles a moderate percentage of outliers without cost blowing up
- Many real problems have high rate of outliers (but sometimes selective choice of random subsets can help)

- **A voting strategy, The Hough transform [Hough, 1959], can handle high percentage of outliers**

- Each point votes separately
- But complexity of search time increases exponentially with the number of model parameters (e.g. 3 for circle)

Summary

- **RANSAC**
 - Algorithm
 - Analysis
 - Number of samples
 - Pros and cons



What we will learn today?

- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector

Some background reading:

Rick Szeliski, Chapter 4.1.1; David Lowe, IJCV 2004

Image matching: a challenging problem



Template

Query



Image matching: a challenging problem



Different photos of the same location, Roma Trevi Fountain


Image matching: a challenging problem

- **Harder case:**



Image matching: a challenging problem


- **Harder Still?**



NASA Mars Rover images

Image matching: a challenging problem

- **Answer Below (Look for tiny colored squares)**





NASA Mars Rover images with SIFT feature matches

Motivation for using local features

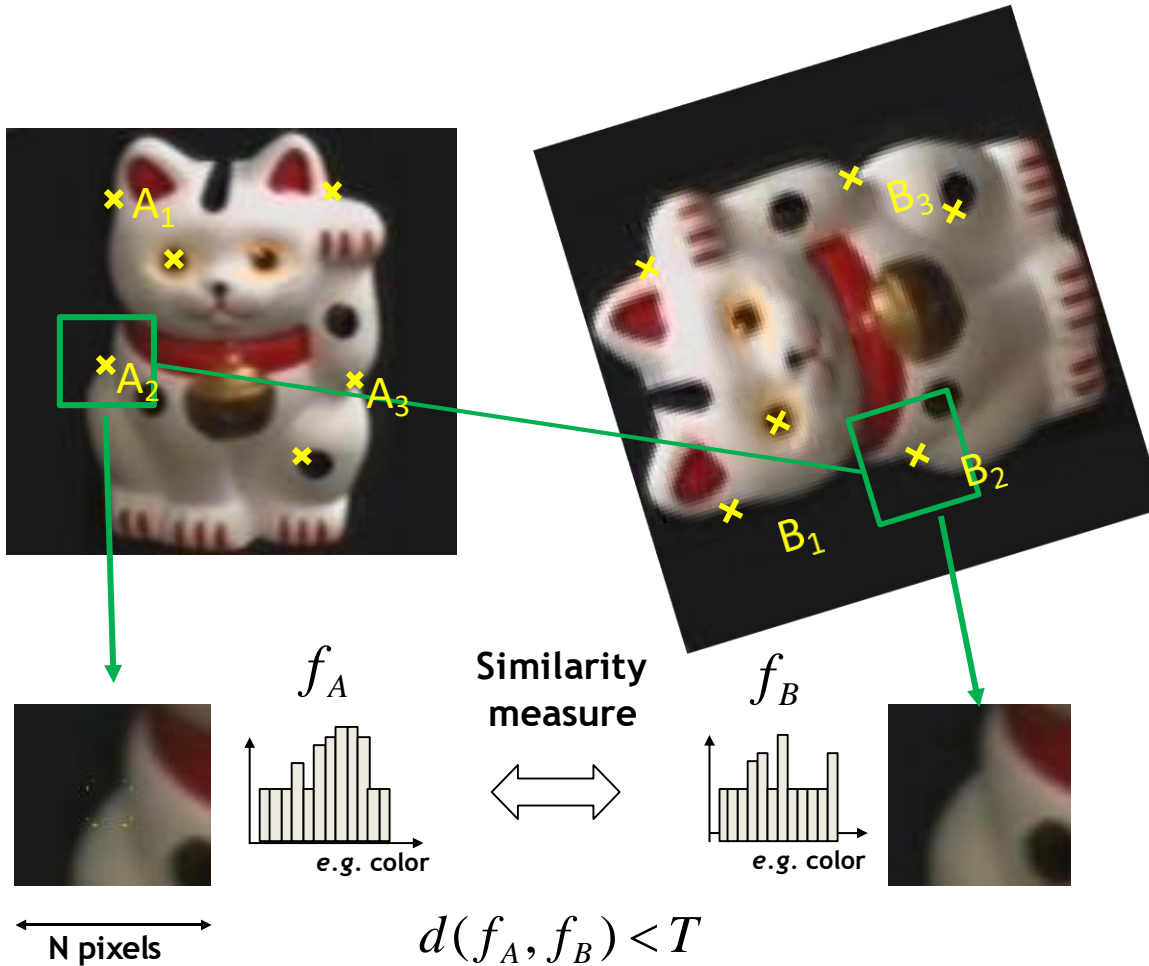
- **Global representations have major limitations.**
- **Instead, describe and match only local regions**
- **Increased robustness to**
 - *Occlusions*
 - *Articulation*



– *Intra-category variations*



General Approach



1. Find a set of distinctive key-points
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor from the normalized region
5. Match local descriptors

Common Requirements

- Problem 1:
 - Detect the same point independently in both images




No chance to match!

We need a repeatable detector!




Common Requirements

- Problem 1:
 - Detect the same point independently in both images
- Problem 2:
 - For each point correctly recognize the corresponding one






We need a repeatable detector!

Feature Invariances: Geometric Transformations



Feature Invariances: Photometric Transformations



- Often modeled as a linear transformation:
 - Scaling + Offset

Requirements for Local Features

- **Region extraction needs to be repeatable and accurate**
 - Invariant to translation, rotation, scale changes
 - Robust or covariant to out-of-plane (\approx affine) transformations
 - Robust to lighting variations, noise, blur, quantization
- **Locality:** Features are local, therefore robust to occlusion and clutter.
- **Quantity:** We need a sufficient number of regions to cover the object.
- **Distinctiveness:** The regions should contain “interesting” structure.
- **Efficiency:** Close to real-time performance.

Many Existing Detectors Available

- Hessian & Harris [Beaudet '78], [Harris '88]
 - Laplacian, DoG [Lindeberg '98], [Lowe '99]
 - Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
 - Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
 - EBR and IBR [Tuytelaars & Van Gool '04]
 - MSER [Matas '02]
 - Salient Regions [Kadir & Brady '01]
 - Others...
-
- Those detectors have become a basic building block for many applications in Computer Vision

Summary

- **Region extraction needs to be repeatable and accurate**
 - Invariant to translation, rotation, scale changes
 - Robust or covariant to out-of-plane (\approx affine) transformations
 - Robust to lighting variations, noise, blur, quantization
- **Local invariant features**
 - Motivation
 - General approach and requirements



What we will learn today?

- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector

Keypoint Localization


- **Goals:**

1. Repeatable detection
2. Precise localization
3. Interesting content


⇒ **Look for two-dimensional signal changes**




What are good keypoints?



(a) Corner

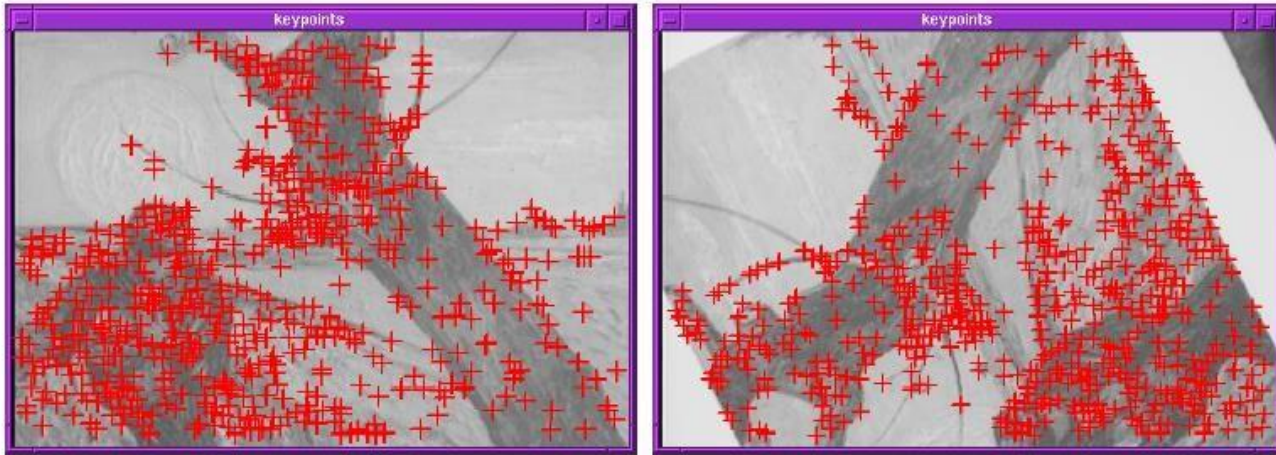


(b) Edge



(c) Textureless

Finding Corners



- **Key property:**


In the region around a corner, the image gradient has two or more dominant directions.

- **Corners are repeatable and distinctive**


Corners as Distinctive Interest Points

- **Design criteria**


1. We should easily recognize the corner point by looking through a small window (**locality**).
2. Shifting the window in any direction should give a large change in intensity (**good localization**).



“flat” region:
no change in
all directions




“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions


Corners versus edges



$$\sum I_x^2 \longrightarrow \text{Large}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$

Corner



$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$


Edge



$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Small}$$

Flat




$$\sum I_x^2 \longrightarrow ??$$

$$\sum I_y^2 \longrightarrow ??$$

Corner

Harris Detector Formulation

- Localize patches that result in large change of intensity when shifted in any direction.
- When we shift by $[u, v]$, the intensity change at the center pixel is:



- *Measure change as intensity difference:*
$$(I(x + u, y + v) - I(x, y))$$
- *That's for a single point, but we have to accumulate over the patch or "small window" around that point...*

“corner”: significant change in all directions

Harris Detector Formulation


- When we shift by $[u, v]$, the change in intensity for the “small window” is:

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Diagram illustrating the Harris Detector Formulation:


- $\sum_{x, y}$: Sum over window (indicated by a green arrow)
- $w(x, y)$: Window function (indicated by an orange arrow)
- $I(x + u, y + v)$: Shifted intensity (indicated by a blue arrow)
- $I(x, y)$: Intensity (indicated by a red arrow)
- $[I(x + u, y + v) - I(x, y)]^2$: Intensity change (indicated by a purple arrow)

Window function $w(x, y) =$



1 in window, 0 outside

or



Gaussian

Harris Detector Formulation

- This measure of change can be approximated by (Taylor expansion):

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

Auto-correlation
matrix of gradients

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Sum over image region – the area we
are checking for corner

*Gradient with
respect to x,
times gradient
with respect to y*

Harris Detector Formulation






Image I



I_x



I_y



$I_x I_y$

- where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Gradient with respect to x , times gradient with respect to y

Sum over image region – the area we are checking for corner


$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$

What Does This Matrix Reveal?

- First, let's consider an axis-aligned corner.
- In that case, the dominant gradient directions align with the x or the y axis.

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

- This means: if either λ is close to 0, then this is not a corner, so look for image windows where both λ are large.
- What if we have a corner that is not aligned with the image axes?




General Case

- Since $M = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$ is symmetric, we can re-rewrite

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R \cdot \textbf{(Eigenvalue decomposition)}$$


- We can think of M as an ellipse with its axis lengths determined by the eigenvalues λ_1 and λ_2 ; and its orientation determined by R



- A rotated corner would produce the same eigenvalues as its non-rotated version.

Interpreting the Eigenvalues

- Classification of image points using eigenvalues of M :




Corner Response Function

$$\theta = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

Fast approximation
[Harris and Stephens,
1988]

- Avoid computing the eigenvalues
- α : constant
(0.04 to 0.06)



Window Function $w(x,y)$


$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Option 1: uniform window

- Sum over square window

$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Problem: not rotation invariant




1 in window, 0 outside

- Option 2: Smooth with Gaussian

- Gaussian already performs weighted sum

$$M = g(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Result is rotation invariant



Gaussian

Summary: Harris Detector


- Compute second moment matrix (autocorrelation matrix)

$$M(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$


σ_D : for Gaussian in the derivative calculation

σ_I : for Gaussian in the windowing function

1. Image derivatives



2. Square of derivatives



3. Gaussian filter $g(\sigma_I)$



4. Cornerness function – two strong eigenvalues

$$\begin{aligned} \theta &= \det[M(\sigma_I, \sigma_D)] - \alpha [\text{trace}(M(\sigma_I, \sigma_D))]^2 \\ &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha [g(I_x^2) + g(I_y^2)]^2 \end{aligned}$$

5. Perform non-maximum suppression




Harris Detector: Workflow

- Input image




Harris Detector: Workflow

- Input Image
- Compute corner response function θ



Harris Detector: Workflow

- Input Image
- Compute corner response function θ
- Take only the local maxima of θ , where $\theta > \text{threshold}$




Harris Detector: Workflow


- Input Image
- Compute corner response function θ
- Take only the local maxima of θ , where $\theta > \text{threshold}$



Harris Detector – Responses



Effect: A very precise corner detector.



Harris Detector – Responses



Harris Detector – Responses



- Results are well suited for finding stereo correspondences




Harris Detector: Properties

- Translation invariance?

Harris Detector: Properties

- Translation invariance
- Rotation invariance?




Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response θ is invariant to image rotation

Harris Detector: Properties

- Translation invariance
- Rotation invariance
- Scale invariance?



Not invariant to image scale!

Summary

- Harris corner detector
 - Formulation
 - Examples